

实验代码



“操作系统原理与实践”实验报告

基于内核栈切换的进程切换

问题1： 由于Linux 0.11进程的内核栈和该进程的PCB在同一页内存上（一块4KB大小的内存），其中PCB位于这页内存的低地址，栈位于这页内存的高地址；加4096就可以得到内核栈地址。tss.ss0是内核数据段，现在只用一个tss，因此不需要设置了。问题2： eax =0 为了与父进程区分开 cpy_process() 让eax=0 这段代码中的ebx和ecx来自copy_process()的形参，是段寄存器，ebp是用户栈地址，一定要设置，不设置子进程就没有用户栈了 问题3： 这两句代码的含义是重新取一下段寄存器fs的值，这两句话必须要加、也必须要出现在切换完LDT之后，这是因为通过fs访问进程的用户态内存，LDT切换完成就意味着切换了分配给进程的用户态内存地址空间，所以前一个fs指向的是上一个进程的用户态内存，而现在需要执行下一个进程的用户态内存，所以就需要用这两条指令来重取fs。出现在LDT之前访问的就还是上一个进程的用户态内存

```
shiyanolou@86d4606c34fe: ~/oslab/oslab/linux-0.11/include/linux
"xchgl %%ecx,current\n\t" \
"ljmp *%0\n\t" \
"cmpl %%ecx,last_task_used_math\n\t" \
"jne 1f\n\t" \
"clts\n\t" \
"1:"
"pushl %%ebp \n\t" \
"movl %%esp,%%ebp \n\t" \
"pushl %%ecx\n\t" \
"pushl %%ebx\n\t" \
"pushl %%eax \n\t" \
"movl 8(%%ebp),%%ebx \n\t"\
"cmpl %%ebx,current\n\t"\
"je 1f\n\t" \
"movl %%ebx,%%eax\n\t" \
"xchgl %%eax,current \n\t" \
"movl tss,%%ecx\n\t" \
"addl $4096,%%ebx \n\t"\
"movl %%ebx,ESP0(%%ecx) \n\t"\
"movl %%esp,KERNEL_STACK(%%eax)\n\t"\
"movl 8(%%ebp),%%ebx\n\t"\
"movl KERNEL_STACK(%%ebx),%%esp\n\t"\
"movl 12(%%ebp),%%ecx\n\t"\
"lldt %%cx\n\t" \
"movl $0x17,%%ecx\n\t"\
"mov %%cx,%%fs\n\t"\
"cmpl %%eax,last_task_used_math\n\t"\
"jne 1f\n\t"\
"clts\n\t"\
"1: popl %%eax\n\t"\
"popl %%ebx\n\t"\
"popl %%ecx\n\t"\
"popl %%ebp\n\t"\
ret
::"m" (*&_tmp.a),"m" (*&_tmp.b), \
"d" (_TSS(n)),"c" ((long) task[n]))]; \

shiyanolou@86d4606c34fe: ~/oslab/oslab/linux-0.11/kernel
/* information (task[nr]) and sets up the necessary registers. It
 * also copies the data segment in it's entirety.
 */
int copy_process(int nr,long ebp,long edi,long esi,long gs,long none,
                long ebx,long ecx,long edx,
                long fs,long es,long ds,
                long eip,long cs,long eflags,long esp,long ss)
{
    struct task_struct *p;
    int i;
    struct file *f;

    p = (struct task_struct *) get_free_page();
    if (!p)
        return -EAGAIN;
    task[nr] = p;
    *p = *current; /* NOTE! this doesn't copy the supervisor stack */
    p->state = TASK_UNINTERRUPTIBLE;
    p->pid = last_pid;
    p->father = current->pid;
    p->counter = p->priority;
    p->signal = 0;
    p->alarm = 0;
    p->leader = 0; /* process leadership doesn't inherit */
    p->utime = p->stime = 0;
    p->cutime = p->cstime = 0;
    p->start_time = jiffies;
    p->tss.back_link = 0;
    p->tss.esp0 = PAGE_SIZE + (long) p;
    *(--krnstack) = ss & 0xffff;
    *(--krnstack) = esp;
    *(--krnstack) = eflags;
    *(--krnstack) = cd & 0xffff;
    *(--krnstack) = eip;
    p->tss.ss0 = 0x10;
    p->tss.eip = eip;
}

shiyanolou@86d4606c34fe: ~/oslab/oslab/linux-0.11/kernel
movl current,%eax # task[0] cannot have signals
cmpl task,%eax
je 3f
cmpw $0x0f,CS(%esp) # was old code segment supervisor ?
jne 3f
cmpw $0x17,OLDSS(%esp) # was stack segment = 0x17 ?
jne 3f
movl signal(%eax),%ebx
movl blocked(%eax),%ecx
notl %ecx
andl %ebx,%ecx
bsfl %ecx,%ecx
je 3f
btrl %ecx,%ebx
movl %ebx,signal(%eax)
incl %ecx
pushl %ecx
call do_signal
popl %eax
3: popl %eax
popl %ebx
popl %ecx
popl %edx
popl %edi
popl %eip
pop %gs
pop %fs
pop %es
pop %ds
iret

.align 2
coprocessor_error:
push %ds
push %es
push %fs
}

shiyanolou@86d4606c34fe: ~/oslab/oslab/linux-0.11/kernel
movl current,%eax # task[0] cannot have signals
cmpl task,%eax
je 3f
cmpw $0x0f,CS(%esp) # was old code segment supervisor ?
jne 3f
cmpw $0x17,OLDSS(%esp) # was stack segment = 0x17 ?
jne 3f
movl signal(%eax),%ebx
movl blocked(%eax),%ecx
notl %ecx
andl %ebx,%ecx
bsfl %ecx,%ecx
je 3f
btrl %ecx,%ebx
movl %ebx,signal(%eax)
incl %ecx
pushl %ecx
call do_signal
popl %eax
3: popl %eax
popl %ebx
popl %ecx
popl %edx
popl %edi
popl %eip
pop %gs
pop %fs
pop %es
pop %ds
iret

.align 2
coprocessor_error:
push %ds
push %es
push %fs
}
```